Kevin Tong - Bird Image Classification

Packages

In [00]:
```
import datetime
import time
from packaging import version
from collections import Counter
import numpy as np
import pandas as pd

import matplotlib as mpl  # EA
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.metrics import confusion_matrix, classification_report
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import mean_squared_error as MSE
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split

import tensorflow as tf
from tensorflow.keras.utils import to_categorical
from tensorflow import keras
from tensorflow.keras import layers, models
from tensorflow.keras.models import Sequential
import tensorflow.keras.backend as z
from tensorflow.keras.utils import plot_model
from tensorflow.keras.layers import Conv2D, MaxPool2D, BatchNormalization, Dropout, Flatten, Dense
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from tensorflow.keras.preprocessing import image
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.layers import Dropout
import os
import cv2
from PIL import Image
from tensorflow.keras.regularizers import l2
```

In [2]:
```
# from google.colab import drive
# drive.mount('/content/drive')
data_dir = '/content/drive/MyDrive/Northwestern MSDS/MSDS 458/458_Final_Notebooks/spectrogram_output'
```

Preprocessing

In [3]:
```
# Initialize a set to store unique classes (bird species)
unique_classes = set()

# Define a function to load and preprocess images
def load_and_preprocess_images(image_paths):
    images = []
    for path in image_paths:
        img = Image.open(path)
        img = img.resize((224, 224))  # Resize to your desired dimensions
        img = img.convert("RGB")  # Convert to RGB (remove alpha channel if present)
        img = np.array(img) / 255.0  # Normalize pixel values
        images.append(img)
    return np.array(images)

# Loop through directories and load images
all_images = []
for bird_folder in os.listdir(data_dir):
    bird_folder_path = os.path.join(data_dir, bird_folder)
    if os.path.isdir(bird_folder_path) and len(os.listdir(bird_folder_path)) > 1:
        # Add the current folder name to the set of unique classes
        unique_classes.add(bird_folder)

        image_paths = [os.path.join(bird_folder_path, img) for img in os.listdir(bird_folder_path)]
        images = load_and_preprocess_images(image_paths)
        all_images.extend(images)

all_images = np.array(all_images)

# Calculate the total number of classes
num_classes = len(unique_classes)
print("Number of classes:", num_classes)
```

```
Number of classes: 110
```

In [4]:
```
# Load and preprocess labels
all_labels = []
for bird_folder in os.listdir(data_dir):
    bird_folder_path = os.path.join(data_dir, bird_folder)
    if os.path.isdir(bird_folder_path) and len(os.listdir(bird_folder_path)) > 1:
        label = bird_folder  # Assuming folder name is the label
        labels = [label] * len(os.listdir(bird_folder_path))
        all_labels.extend(labels)

# Convert labels to numerical format (e.g., one-hot encoding)
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

label_encoder = LabelEncoder()
integer_encoded = label_encoder.fit_transform(all_labels)
onehot_encoder = OneHotEncoder(sparse=False)
encoded_labels = onehot_encoder.fit_transform(integer_encoded.reshape(-1, 1))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
  warnings.warn(
```

In [5]:
```
from sklearn.model_selection import train_test_split

# Split data into training and temporary set (which includes validation and test data)
train_images, temp_images, train_labels, temp_labels = train_test_split(all_images, encoded_labels, test_size=0.2, random_state=42)

# Split temporary set into validation and test sets
val_images, test_images, val_labels, test_labels = train_test_split(temp_images, temp_labels, test_size=0.5, random_state=42)
```

In [6]:
```
print("Train images shape:", train_images.shape)
print("Validation images shape:", val_images.shape)
print("Test images shape:", test_images.shape)
```

```
Train images shape: (1725, 224, 224, 3)
Validation images shape: (216, 224, 224, 3)
Test images shape: (216, 224, 224, 3)
```

Neural Network

In [103]:
```
start_time = time.time()

# Define a complex CNN model
model = models.Sequential([
    # Convolutional layers
    layers.Conv2D(64, (3, 3), activation='relu', input_shape=(224, 224, 3)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(256, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(512, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),

    # Fully connected layers with dropout and batch normalization
    layers.Dense(512, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Dense(256, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    layers.Dense(128, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.5),

    # Output layer
    layers.Dense(num_classes, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

In [104]:
```
# Print the model summary
model.summary()
```

```
Model: "sequential_16"
_________________________________________________________________
 Layer (type)                Output Shape              Param #
=================================================================
 conv2d_35 (Conv2D)          (None, 222, 222, 64)      1792

 max_pooling2d_35 (MaxPoolin  (None, 111, 111, 64)     0
 g2D)

 conv2d_36 (Conv2D)          (None, 109, 109, 128)     73856

 max_pooling2d_36 (MaxPoolin  (None, 54, 54, 128)      0
 g2D)

 conv2d_37 (Conv2D)          (None, 52, 52, 256)       295168

 max_pooling2d_37 (MaxPoolin  (None, 26, 26, 256)      0
 g2D)

 conv2d_38 (Conv2D)          (None, 24, 24, 512)       1180160

 max_pooling2d_38 (MaxPoolin  (None, 12, 12, 512)      0
 g2D)

 flatten_16 (Flatten)        (None, 73728)             0

 dense_38 (Dense)            (None, 512)               37749248

 batch_normalization (BatchN  (None, 512)              2048
 ormalization)

 dropout_26 (Dropout)        (None, 512)               0

 dense_39 (Dense)            (None, 256)               131328

 batch_normalization_1 (Batc  (None, 256)              1024
 hNormalization)

 dropout_27 (Dropout)        (None, 256)               0

 dense_40 (Dense)            (None, 128)               32896

 batch_normalization_2 (Batc  (None, 128)              512
 hNormalization)

 dropout_28 (Dropout)        (None, 128)               0

 dense_41 (Dense)            (None, 110)               14190

=================================================================
Total params: 39,482,222
Trainable params: 39,480,430
Non-trainable params: 1,792
_________________________________________________________________
```

In [105]:
```
# Train the model
history = model.fit(train_images, train_labels, epochs=200, batch_size=64,
                    validation_data=(val_images, val_labels),callbacks=[
                    tf.keras.callbacks.ModelCheckpoint("CNN_model.h5",save_best_only=True,save_weights_only=False)
                    ,tf.keras.callbacks.EarlyStopping(monitor='val_accuracy', patience=3),
                    ])

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print('Test accuracy:', test_acc)
```

```
Epoch 1/200
27/27 [==============================] - 11s 195ms/step - loss: 5.4769 - accuracy: 0.0122 - val_loss: 24.4894 - val_accuracy: 0.0046
Epoch 2/200
27/27 [==============================] - 2s 76ms/step - loss: 5.1489 - accuracy: 0.0214 - val_loss: 13.1227 - val_accuracy: 0.0093
Epoch 3/200
27/27 [==============================] - 2s 77ms/step - loss: 4.8725 - accuracy: 0.0388 - val_loss: 6.3260 - val_accuracy: 0.0139
Epoch 4/200
27/27 [==============================] - 3s 112ms/step - loss: 4.6810 - accuracy: 0.0510 - val_loss: 4.9028 - val_accuracy: 0.0185
Epoch 5/200
27/27 [==============================] - 2s 71ms/step - loss: 4.4838 - accuracy: 0.0620 - val_loss: 4.8687 - val_accuracy: 0.0278
Epoch 6/200
27/27 [==============================] - 1s 43ms/step - loss: 4.2667 - accuracy: 0.0852 - val_loss: 6.0822 - val_accuracy: 0.0370
Epoch 7/200
27/27 [==============================] - 2s 91ms/step - loss: 4.0887 - accuracy: 0.1038 - val_loss: 4.6150 - val_accuracy: 0.0880
Epoch 8/200
27/27 [==============================] - 1s 43ms/step - loss: 4.0152 - accuracy: 0.1113 - val_loss: 16.6345 - val_accuracy: 0.0046
Epoch 9/200
27/27 [==============================] - 1s 43ms/step - loss: 3.8790 - accuracy: 0.1299 - val_loss: 7.7089 - val_accuracy: 0.0278
Epoch 10/200
27/27 [==============================] - 1s 43ms/step - loss: 3.7315 - accuracy: 0.1357 - val_loss: 7.7624 - val_accuracy: 0.0417
7/7 [==============================] - 0s 11ms/step - loss: 7.5784 - accuracy: 0.0278
Test accuracy: 0.027777778473835
```

In [106]:
```
end_time = time.time()
runtime = end_time - start_time
print('Total runtime:', runtime, 'seconds')
```

```
Total runtime: 34.619449853897095 seconds
```

In [107]:
```
# Plot training and validation accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()
plt.show()
```

In [108]:
```
# Plot test accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(len(history.history['val_accuracy']), test_acc, 'ro', label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Test Accuracy')
plt.legend()
plt.show()
```